

L02- Converting and Deploying AI Models using TensorFlow Lite

Option B - Practical

Part 1: Setting Up the Development Environment

Step 1: Verify Python and TensorFlow Installation

```
In [3]: !python --version
```

Python 3.12.7

```
In [4]: !pip show tensorflow
```

Name: tensorflow
Version: 2.18.0
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: <https://www.tensorflow.org/>
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: C:\Users\banke\anaconda3\Lib\site-packages
Requires: tensorflow-intel
Required-by:

```
In [5]: !pip install tensorflow
```

Requirement already satisfied: tensorflow in c:\users\banke\anaconda3\lib\site-packages (2.18.0)

Requirement already satisfied: tensorflow-intel==2.18.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow) (2.18.0)

Requirement already satisfied: absl-py>=1.0.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (2.1.0)

Requirement already satisfied: astunparse>=1.6.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=24.3.25 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (25.1.24)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (0.6.0)

Requirement already satisfied: google-pasta>=0.1.1 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (0.2.0)

Requirement already satisfied: libclang>=13.0.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (18.1.1)

Requirement already satisfied: opt-einsum>=2.3.2 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (3.4.0)

Requirement already satisfied: packaging in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (24.1)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (4.25.3)

Requirement already satisfied: requests<3,>=2.21.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (2.32.3)

Requirement already satisfied: setuptools in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (75.1.0)

Requirement already satisfied: six>=1.12.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (1.16.0)

Requirement already satisfied: termcolor>=1.1.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (2.5.0)

Requirement already satisfied: typing-extensions>=3.6.6 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (4.11.0)

Requirement already satisfied: wrapt>=1.11.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (1.14.1)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (1.70.0)

Requirement already satisfied: tensorboard<2.19,>=2.18 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (2.18.0)

Requirement already satisfied: keras>=3.5.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (3.8.0)

Requirement already satisfied: numpy<2.1.0,>=1.26.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (1.26.4)

Requirement already satisfied: h5py>=3.11.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (3.11.0)

Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorflow-intel==2.18.0->tensorflow) (0.4.1)

Requirement already satisfied: wheel<1.0,>=0.23.0 in c:\users\banke\anaconda3\lib\site-packages (from astunparse>=1.6.0->tensorflow-intel==2.18.0->tensorflow) (0.44.0)

Requirement already satisfied: rich in c:\users\banke\anaconda3\lib\site-packages (from keras>=3.5.0->tensorflow-intel==2.18.0->tensorflow) (13.7.1)

Requirement already satisfied: namex in c:\users\banke\anaconda3\lib\site-packages (from keras>=3.5.0->tensorflow-intel==2.18.0->tensorflow) (0.0.8)

Requirement already satisfied: optree in c:\users\banke\anaconda3\lib\site-packages (from keras>=3.5.0->tensorflow-intel==2.18.0->tensorflow) (0.14.0)

Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\banke\anaconda3

\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.18.0->tensorflow) (3.3.2)
 Requirement already satisfied: idna<4,>=2.5 in c:\users\banke\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.18.0->tensorflow) (3.7)
 Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\banke\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.18.0->tensorflow) (2.2.3)
 Requirement already satisfied: certifi>=2017.4.17 in c:\users\banke\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorflow-intel==2.18.0->tensorflow) (2024.12.14)
 Requirement already satisfied: markdown>=2.6.8 in c:\users\banke\anaconda3\lib\site-packages (from tensorboard<2.19,>=2.18->tensorflow-intel==2.18.0->tensorflow) (3.4.1)
 Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in c:\users\banke\anaconda3\lib\site-packages (from tensorboard<2.19,>=2.18->tensorflow-intel==2.18.0->tensorflow) (0.7.2)
 Requirement already satisfied: werkzeug>=1.0.1 in c:\users\banke\anaconda3\lib\site-packages (from tensorboard<2.19,>=2.18->tensorflow-intel==2.18.0->tensorflow) (3.0.3)
 Requirement already satisfied: MarkupSafe>=2.1.1 in c:\users\banke\anaconda3\lib\site-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow-intel==2.18.0->tensorflow) (2.1.3)
 Requirement already satisfied: markdown-it-py>=2.2.0 in c:\users\banke\anaconda3\lib\site-packages (from rich->keras>=3.5.0->tensorflow-intel==2.18.0->tensorflow) (2.2.0)
 Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\banke\anaconda3\lib\site-packages (from rich->keras>=3.5.0->tensorflow-intel==2.18.0->tensorflow) (2.15.1)
 Requirement already satisfied: mdurl~=0.1 in c:\users\banke\anaconda3\lib\site-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow-intel==2.18.0->tensorflow) (0.1.0)

Part 2: Creating and Training an AI Model

Step 3: Load the MNIST Dataset

```

In [7]: import tensorflow as tf
        from tensorflow import keras
        from tensorflow.keras.datasets import mnist
        import numpy as np
        import matplotlib.pyplot as plt
  
```


```

In [8]: # Load MNIST dataset
        (x_train, y_train), (x_test, y_test) = mnist.load_data()

        # Normalize data (scale pixel values between 0 and 1)
        x_train, x_test = x_train / 255.0, x_test / 255.0
  
```

```

In [9]: # Show sample images
        plt.figure(figsize=(5,5))
        for i in range(9): plt.subplot(3,3,i+1)
        plt.imshow(x_train[i], cmap="gray")
        plt.axis('off')
        plt.show()
  
```



Step 4: Define and Train a Neural Network

Create a simple feedforward neural network using Keras:

```
In [11]: # Define model architecture
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(), # Input Layer
    tf.keras.layers.Dense(128, activation='relu'), # Hidden Layer
    tf.keras.layers.Dense(10, activation='softmax') # Output layer (10 classes)
])

In [12]: # Compile model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['a

In [13]: # Train model
model.fit(x_train, y_train, epochs=5, validation_data=(x_test, y_test))
```

```
Epoch 1/5
1875/1875 ————— 4s 1ms/step - accuracy: 0.8757 - loss: 0.4398 - val_a
ccuracy: 0.9590 - val_loss: 0.1415
Epoch 2/5
1875/1875 ————— 3s 2ms/step - accuracy: 0.9642 - loss: 0.1233 - val_a
ccuracy: 0.9699 - val_loss: 0.1016
Epoch 3/5
1875/1875 ————— 3s 2ms/step - accuracy: 0.9773 - loss: 0.0777 - val_a
ccuracy: 0.9718 - val_loss: 0.0929
Epoch 4/5
1875/1875 ————— 3s 1ms/step - accuracy: 0.9824 - loss: 0.0576 - val_a
ccuracy: 0.9758 - val_loss: 0.0793
Epoch 5/5
1875/1875 ————— 3s 1ms/step - accuracy: 0.9873 - loss: 0.0442 - val_a
ccuracy: 0.9745 - val_loss: 0.0826
```

```
Out[13]: <keras.src.callbacks.history.History at 0x1a6023de720>
```

```
In [14]: # Save trained model
model.save("mnist_model.keras", include_optimizer=False)
```

```
In [15]: print("Model training complete and saved as mnist_model.keras")
```

Model training complete and saved as mnist_model.keras

Part 3: Converting and Saving the Model

Step 5: Convert the Model to TensorFlow Lite Format

```
In [17]: # Load trained model
model = tf.keras.models.load_model("mnist_model.keras")
```

```
In [18]: # Convert to TensorFlow Lite
converter = tf.lite.TFLiteConverter.from_keras_model(model)
tflite_model = converter.convert()
```

```
INFO:tensorflow:Assets written to: C:\Users\banke\AppData\Local\Temp\tmpsyle83q\ass
ets
```

```
INFO:tensorflow:Assets written to: C:\Users\banke\AppData\Local\Temp\tmpsyle83q\ass
ets
```

Saved artifact at 'C:\Users\banke\AppData\Local\Temp\tmpsyle83q'. The following end points are available:

* Endpoint 'serve'

```
args_0 (POSITIONAL_ONLY): TensorSpec(shape=(32, 28, 28), dtype=tf.float32, name='i
nput_layer')
```

Output Type:

```
TensorSpec(shape=(32, 10), dtype=tf.float32, name=None)
```

Captures:

```
1812514052944: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
1812514047952: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
1812514048528: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
1812514051600: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
In [19]: # Save the converted model
with open("mnist_model.tflite", "wb") as f:
    f.write(tflite_model)
```

```
In [20]: print("Model successfully converted and saved as mnist_model.tflite")
```

Model successfully converted and saved as mnist_model.tflite

Part 4: Loading and Running Inference with TensorFlow Lite

Step 6: Load the Converted Model Using TensorFlow Lite Interpreter

```
In [22]: # Load TensorFlow Lite model
interpreter = tf.lite.Interpreter(model_path="mnist_model.tflite")
interpreter.allocate_tensors()
```

```
In [23]: # Get input and output tensor details
input_details = interpreter.get_input_details()
output_details = interpreter.get_output_details()
```

```
In [24]: print("Input Details:", input_details)
```

Input Details: [{'name': 'serving_default_input_layer:0', 'index': 0, 'shape': array([32, 28, 28]), 'shape_signature': array([32, 28, 28]), 'dtype': <class 'numpy.float32'>, 'quantization': (0.0, 0), 'quantization_parameters': {'scales': array([], dtype=float32), 'zero_points': array([], dtype=int32), 'quantized_dimension': 0}, 'sparsity_parameters': {}}]

```
In [25]: print("Output Details:", output_details)
```

Output Details: [{'name': 'StatefulPartitionedCall_1:0', 'index': 9, 'shape': array([32, 10]), 'shape_signature': array([32, 10]), 'dtype': <class 'numpy.float32'>, 'quantization': (0.0, 0), 'quantization_parameters': {'scales': array([], dtype=float32), 'zero_points': array([], dtype=int32), 'quantized_dimension': 0}, 'sparsity_parameters': {}}]

Step 7: Perform Inference with TensorFlow Lite

```
In [27]: # Select a test image
test_image = x_test[0].astype(np.float32)
```

```
In [28]: # Ensure data type matches model input
test_image = np.expand_dims(test_image, axis=0) # Add batch dimension (1, 28, 28)
test_image = np.tile(test_image, (32, 1, 1)) # (32, 28, 28)
# Convert to float32 (required for TensorFlow Lite)
test_image = test_image.astype(np.float32)
```

```
In [29]: print("Expected input shape:", input_details[0]['shape'])
```


Expected input shape: [32 28 28]

```
In [30]: # Set the input tensor
interpreter.set_tensor(input_details[0]['index'], test_image)
```

```
In [31]: # Run inference
interpreter.invoke()
```

```
In [32]: # Get the prediction
output_data = interpreter.get_tensor(output_details[0]['index'])
predicted_label = np.argmax(output_data)
```

```
In [33]: # Display the image and prediction
plt.imshow(x_test[0], cmap="gray")
plt.title(f"Predicted Label: {predicted_label}, Actual Label: {y_test[0]}")
plt.show()
```



```
In [ ]:
```