

Adebanke Damola-Fashola

ITAI 3377- AI at the Edge and IIOT Environments

Professor Patricia McManus

February 11, 2025.

Simulation Documentation

Deploying a Simple AI Model on a Simulated Edge Device using Visual Studio Code

Step 1-Set Up the Environment:

- **Install Python:** I updated the Python version on my system to the latest version 3.13.2.
- **Install VS Code:** I downloaded Visual Studio Code and installed it on my existing Anaconda platform.
- **Install TensorFlow:** I ran code 'pip install tensorflow numpy matplotlib' on the VS code terminal.
- **Install Edge Impulse CLI:** I downloaded Node.js from <https://nodejs.org/en> and installed it on my system. Then I ran code 'npm install -g edge-impulse-cli' on the powershell terminal.

Step 2 - Prepare the Dataset:

- **Load and Preprocess the Data:** I created a new file named deploy_mnist and selected All Files as the type. This file was used for coding in Python. I loaded the MNIST dataset.
- **Input:**

```
Welcome  deploy_mnist.py X
deploy_mnist.py > ...
1  import tensorflow as tf
2  from tensorflow.keras.datasets import mnist # type: ignore
3
4  # Load dataset
5  (x_train, y_train), (x_test, y_test) = mnist.load_data()
6
7  # Normalize the pixel values (0-1)
8  x_train, x_test = x_train / 255.0, x_test / 255.0
9
10 # Reshape to match the input shape for CNN (28x28x1)
11 x_train = x_train.reshape((-1, 28, 28, 1))
12 x_test = x_test.reshape((-1, 28, 28, 1))
13
14 # Print dataset shape
15 print(f"Train shape: {x_train.shape}, Test shape: {x_test.shape}")
16
```

- **Output:**

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS
PS C:\Users\banke\Desktop\ITAI 3377\Module 03\New folder> & C:/Users/banke/anaconda3/python.exe "c:/Users/banke/Desktop/ITAI 3377/Module 03/New folder/deploy_mnist.py"
2025-02-14 20:33:49.385876: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-02-14 20:33:55.136925: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
Train shape: (60000, 28, 28, 1), Test shape: (10000, 28, 28, 1)
```

Step 3 – Train a Simple AI Model:

- **Define and Compile the Model:**
- *Input:*

```

17 # Define a simple CNN model
18 model = tf.keras.Sequential([
19     tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
20     tf.keras.layers.MaxPooling2D((2, 2)),
21     tf.keras.layers.Flatten(),
22     tf.keras.layers.Dense(128, activation='relu'),
23     tf.keras.layers.Dense(10, activation='softmax')
24 ])
25
26 # Compile the model
27 model.compile(optimizer='adam',
28               loss='sparse_categorical_crossentropy',
29               metrics=['accuracy'])
30
31 # Print model summary
32 model.summary()

```

- *Output:*

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 128)	692,352
dense_1 (Dense)	(None, 10)	1,290

Total params: 693,962 (2.65 MB)
 Trainable params: 693,962 (2.65 MB)
 Non-trainable params: 0 (0.00 B)

- **Train the Model:**






- *Input:*

```

35 # Train the model
36 model.fit(x_train, y_train, epochs=5, validation_data=(x_test, y_test))
37
38 # Save the trained model
39 model.save("my_model.keras")
40 print("Model training complete and saved as my_model.keras")

```

- *Output:*

```
Epoch 1/5
1875/1875  12s 6ms/step - accuracy: 0.9130 - loss: 0.2909 - val_accuracy: 0.9796
- val_loss: 0.0601
Epoch 2/5
1875/1875  11s 6ms/step - accuracy: 0.9833 - loss: 0.0545 - val_accuracy: 0.9835
- val_loss: 0.0459
Epoch 3/5
1875/1875  10s 5ms/step - accuracy: 0.9912 - loss: 0.0311 - val_accuracy: 0.9867
- val_loss: 0.0405
Epoch 4/5
1875/1875  10s 5ms/step - accuracy: 0.9943 - loss: 0.0181 - val_accuracy: 0.9869
- val_loss: 0.0411
Epoch 5/5
1875/1875  10s 5ms/step - accuracy: 0.9955 - loss: 0.0135 - val_accuracy: 0.9862
- val_loss: 0.0466
Model training complete and saved as my_model.keras
```

Step 4 - Convert and Deploy the Model:

- Convert the Model to TFLite:

- *Input:*

```
42 # Convert the trained model to TensorFlow Lite format
43 converter = tf.lite.TFLiteConverter.from_keras_model(model)
44 tflite_model = converter.convert()
45
46 # Save the TFLite model
47 with open("model.tflite", "wb") as f:
48     f.write(tflite_model)
49
50 print("Model converted to TensorFlow Lite format and saved as model.tflite")
```

- *Output:*

Saved artifact at 'C:\Users\banke\AppData\Local\Temp\tmpgykelhs1'. The following endpoints are available:

```
* Endpoint 'serve'
```

```
args_0 (POSITIONAL_ONLY): TensorSpec(shape=(None, 28, 28, 1), dtype=tf.float32, name='keras_tensor')
```

Output Type:

```
TensorSpec(shape=(None, 10), dtype=tf.float32, name=None)
```

Captures:

```
2361092020688: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
2361092021648: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
2361092022992: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
2361092021840: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
2361092022032: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

```
2361092023568: TensorSpec(shape=(), dtype=tf.resource, name=None)
```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```
W0000 00:00:1739587894.768403 45244 tf_tfl_flatbuffer_helpers.cc:365] Ignored output_format.
```

```
W0000 00:00:1739587894.769467 45244 tf_tfl_flatbuffer_helpers.cc:368] Ignored drop_control_dependency.
```

```
2025-02-14 20:51:34.773673: I tensorflow/cc/saved_model/reader.cc:83] Reading SavedModel from: C:\Users\banke\AppData\Local\Temp\tmpgykelhs1
```

```
2025-02-14 20:51:34.777856: I tensorflow/cc/saved_model/reader.cc:52] Reading meta graph with tags { serve }
```

```
2025-02-14 20:51:34.778064: I tensorflow/cc/saved_model/reader.cc:147] Reading SavedModel debug info (if present) from: C:\Users\banke\AppData\Local\Temp\tmpgykelhs1
```

```
I0000 00:00:1739587894.786496 45244 mlir_graph_optimization_pass.cc:401] MLIR V1 optimization pass is not enabled
```

```
2025-02-14 20:51:34.788322: I tensorflow/cc/saved_model/loader.cc:236] Restoring SavedModel bundle.
```

```
2025-02-14 20:51:34.856389: I tensorflow/cc/saved_model/loader.cc:220] Running initialization op on SavedModel bundle at path: C:\Users\banke\AppData\Local\Temp\tmpgkylhs1
```

```
2025-02-14 20:51:34.870901: I tensorflow/cc/saved_model/loader.cc:466] SavedModel load for tags { serve }; Status: success: OK. Took 97819 microseconds.
```

```
2025-02-14 20:51:34.958644: I tensorflow/compiler/mlir/tensorflo
```

Model converted to TensorFlow Lite format and saved as model.tflite

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

powershell + v [icon] ... ^

```
PS C:\Users\banke\Desktop\ITAI 3377\Module 03\New folder> npm install -g edge-impulse-cli
```

```
npm warn deprecated osenv@0.1.5: This package is no longer supported.
```

```
npm warn deprecated inflight@1.0.6: This module is not supported, and leaks memory. Do not use it. Check out lru-cache if you want a good and tested way to coalesce async requests by a key value, which is much more comprehensive and powerful.
```

```
npm warn deprecated @zeit/dockerignore@0.0.5: "@zeit/dockerignore" is no longer maintained
```

```
npm warn deprecated move-concurrently@1.0.1: This package is no longer supported.
```

```
npm warn deprecated rimraf@2.7.1: Rimraf versions prior to v4 are no longer supported
```

```
npm warn deprecated npmlog@4.1.2: This package is no longer supported.
```

```
npm warn deprecated request-promise@4.2.4: request-promise has been deprecated because it extends the now deprecated request package,
see https://github.com/request/request/issues/3142
```

```
npm warn deprecated glob@7.2.3: Glob versions prior to v9 are no longer supported
```

```
npm warn deprecated are-we-there-yet@1.1.7: This package is no longer supported.
```

```
npm warn deprecated debug@4.1.1: Debug versions >=3.2.0 <3.2.7 || >=4 <4.3.1 have a low-severity ReDoS regression when used in a Node.js environment. It is recommended you upgrade to 3.2.7 or 4.3.1. (https://github.com/visionmedia/debug/issues/797)
```

```
npm warn deprecated debug@4.1.1: Debug versions >=3.2.0 <3.2.7 || >=4 <4.3.1 have a low-severity ReDoS regression when used in a Node .js environment. It is recommended you upgrade to 3.2.7 or 4.3.1. (https://github.com/visionmedia/debug/issues/797)
```

```
npm warn deprecated debug@4.1.1: Debug versions >=3.2.0 <3.2.7 || >=4 <4.3.1 have a low-severity ReDoS regression when used in a Node.js environment. It is recommended you upgrade to 3.2.7 or 4.3.1. (https://github.com/visionmedia/debug/issues/797)
```

```
npm warn deprecated fs-write-stream-atomic@1.0.10: This package is no longer supported.
```

```
npm warn deprecated @xmldom/xmldom@0.8.8: this version has critical issues, please update to the latest version
```

```
npm warn deprecated gauge@2.7.4: This package is no longer supported.
```

```
npm warn deprecated uuid@3.4.0: Please upgrade to version 7 or higher. Older versions may use Math.random() in certain circumstances, which is known to be problematic. See https://v8.dev/blog/math-random for details.
```

```
npm warn deprecated request@2.88.0: request has been deprecated, see https://github.com/request/request/issues/3142
```

changed 558 packages in 34s

49 packages are looking for funding

```
run `npm fund` for details
```

```
PS C:\Users\banke\Desktop\ITAI 3377\Module 03\New folder> edge-impulse-uploader --api-key ei_ce852fd77e5bf0a36eb2c438fa97f5135af6fdef
e4e2cf81bfebd2ee92649a49 model.tflite
Edge Impulse uploader v1.30.4
Endpoints:
  API:      https://studio.edgeimpulse.com
  Ingestion: https://ingestion.edgeimpulse.com

Upload configuration:
  Label:      Not set, will be inferred from file name
  Category:   training
Cannot handle this file, only .wav, .cbor, .json, .jpg, .jpeg, .png, .csv, .txt, .mp4, .avi supported: model.tflite
```

- **Upload the dataset to Edge Impulse:**

Upload data

×

You can upload CBOR, JSON, CSV, Parquet, WAV, JPG, PNG, AVI or MP4 files. You can also upload an annotation file named "info.labels" with your data to assign bounding boxes, labels, and/or metadata. View [Uploader docs](#) to learn more. Alternatively, you can use our [Python SDK](#) to programmatically ingest data in various formats, such as pandas or numpy.

For CSV and Parquet files, [configure the CSV Wizard](#) to define how your files should be processed before uploading files.

Upload mode

- ☒ Select individual files ②
- ☐ Select a folder ②

Select files

[Choose Files](#) 5729 files

Upload into category

- ☒ Automatically split between training and testing ②
- ☐ Training
- ☐ Testing

Label

- ☒ Infer from filename ②
- ☐ Leave data unlabeled ②
- ☐ Enter label:

Upload output

```
[1376/5729] Uploading cat.1375.jpg OK
[1377/5729] Uploading cat.1376.jpg OK
[1378/5729] Uploading cat.1377.jpg OK
[1379/5729] Uploading cat.1380.jpg OK
[1380/5729] Uploading cat.1383.jpg OK
[1381/5729] Uploading cat.1378.jpg OK
[1382/5729] Uploading cat.1381.jpg OK
[1383/5729] Uploading cat.1379.jpg OK
[1384/5729] Uploading cat.1382.jpg OK
[1385/5729] Uploading cat.1385.jpg OK
[1386/5729] Uploading cat.1388.jpg OK
[1387/5729] Uploading cat.1386.jpg OK
[1388/5729] Uploading cat.1389.jpg OK
[1389/5729] Uploading cat.1392.jpg OK
[1390/5729] Uploading cat.1390.jpg OK
[1391/5729] Uploading cat.1387.jpg OK
[1392/5729] Uploading cat.1384.jpg OK
[1393/5729] Uploading cat.1393.jpg OK
[1394/5729] Uploading cat.1391.jpg OK
[1395/5729] Uploading cat.1395.jpg OK
[1396/5729] Uploading cat.1394.jpg OK
```

- **Upload the model to Edge Impulse:**

Upload pretrained model - Step 1: Upload a model

Upload a TensorFlow SavedModel (`.saved_model.zip`), ONNX model (`.onnx`), TensorFlow Lite model (`.tflite`), LGBM model (`.txt`), XGBoost model (`.json`) or pickle model (`.pk1`) to get started.

2. Model performance

Do you want performance characteristics (latency, RAM and ROM) for a specific device?

☐ No, show me performance for a range of device types.

☒ Yes, run performance profiling for: Cortex-M4F 80MHz

Upload progress

Creating job... OK (ID: 29079277)

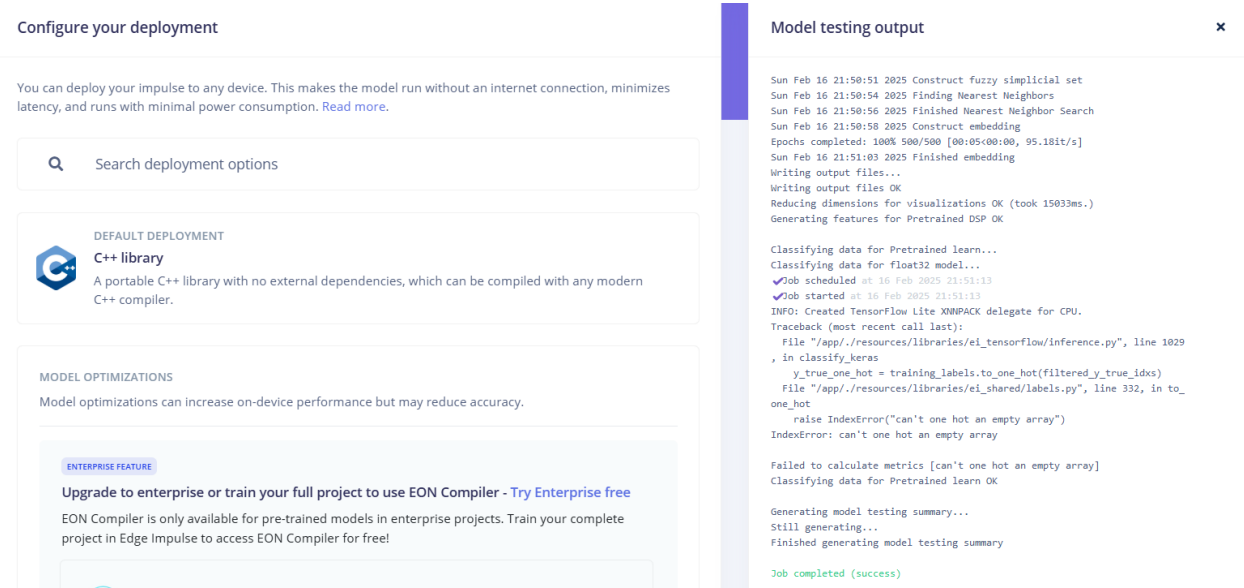
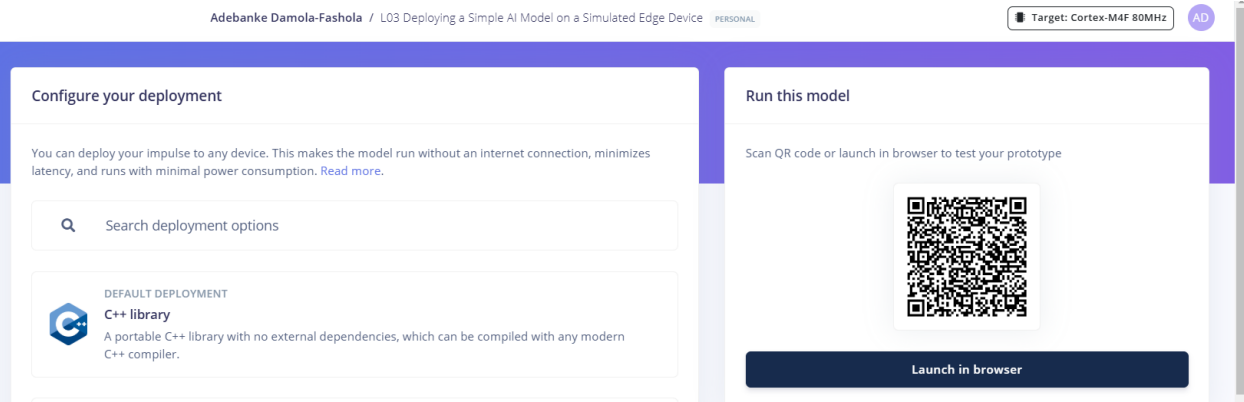
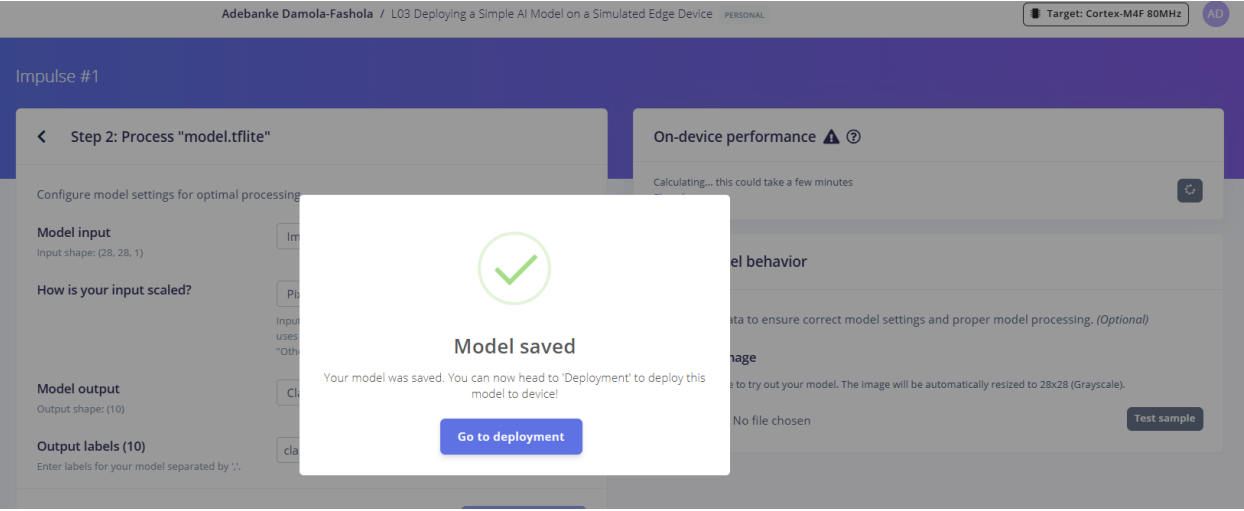
✓Job scheduled at 16 Feb 2025 21:47:49

✓ Job started at 16 Feb 2025 21:47:51

Reading model info...

Reading model info OK

Job completed (success)



- **Retrain the model on Edge Impulse:**

Adebanke Damola-Fashola / L03 Deploying a Simple AI Model on a Simulated Edge Device

Target: Cortex-M4F 80MHz

AD

Retrain model with known parameters

Pretrained DSP

Pretrained learn

Model testing

Train model

Adebanke Damola-Fashola / L03 Deploying a Simple AI Model on a Simulated Edge Device

Target: Cortex-M4F 80MHz

AD

Retrain model with known parameters

✓ Pretrained DSP

✓ Pretrained learn

✓ Model testing

Train model

Build output

(0)

Classifying data for Pretrained learn OK

Generating model testing summary...

Still generating...

Finished generating model testing summary

Running model testing OK

Job completed (success)

- **Use live classification to test:**

Classification result

Summary

Model version: Unoptimized (float32)

Name: dog.12496

Label: dog


CATEGORY	COUNT
cat	0
dog	1
uncertain	0

Detailed result ☐ Show only unknowns

CAT	DOG
0.03	0.97

RAW DATA

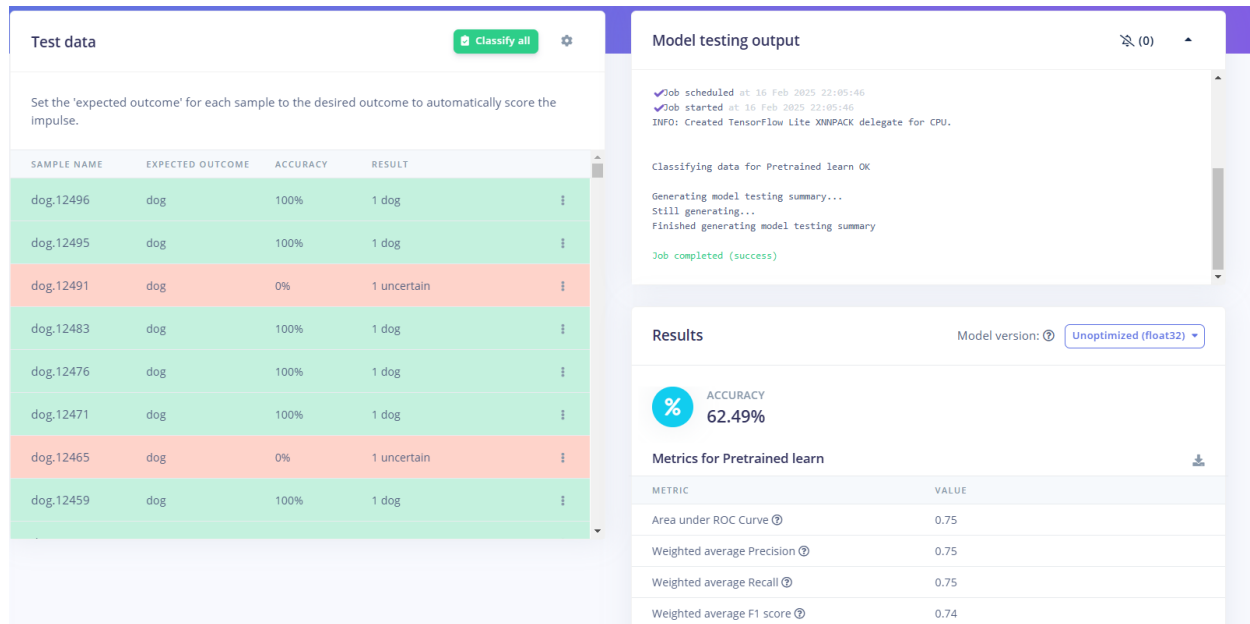
dog.12496



Raw features

0x36312e, 0x3b322e, 0x3a3637, 0x5c696e, 0xd7e4ea, 0xffffffff, 0xffffffff, 0xa8afb0, 0x839091, 0xffff...

Step 5 - Test and Validate the Model:



Below is the breakdown of the key insights of a classification model for distinguishing between cats and dogs on Edge Impulse:

- **Model Performance Overview:**

Accuracy - 62.49%: The model correctly classified 62.49% of test samples.

ROC-AUC Score - 0.75: It suggests a moderately strong ability to distinguish between classes.

Weighted Precision: 0.75

Weighted Recall: 0.75

Weighted F1 Score: 0.74

Precision, Recall, and F1 Score: These metrics indicate a decent balance between precision and recall, though they suggest some misclassification issues.

Metrics for Pretrained learn




METRIC	VALUE
Area under ROC Curve ?	0.75
Weighted average Precision ?	0.75
Weighted average Recall ?	0.75
Weighted average F1 score ?	0.74

Confusion matrix

	CAT	DOG	UNCERTAIN
CAT	56.9%	20.6%	22.5%
DOG	12.2%	68.2%	19.7%
F1 SCORE	0.67	0.72	

Feature explorer ?



- **Confusion Matrix Insights:**

Cat Classification: 56.9% correctly classified as cats, 20.6% misclassified as dogs, and 22.5% categorized as 'uncertain.'

Dog Classification: 68.2% were correctly classified as dogs, 12.2% misclassified as cats, and 19.7% categorized as "uncertain."

F1 Scores: Cats: 0.67 Dogs: 0.72

The model performs slightly better on dogs than on cats.

- **Feature Explorer:**

The scatter plot visualizes feature space clustering. Green (correct classifications) and red (incorrect classifications) are mixed, suggesting overlapping features and contributing to classification challenges.

- **Key Takeaways and Areas for Improvement:**

The model's accuracy is moderate (62.49%). A high uncertainty rate (19-22%) suggests that the model struggles with certain cases. Dogs are classified better than cats based on F1 scores.

Performance can be improved using a more balanced dataset with clean and accurately labeled data during training. Additionally, exploring hyperparameter tuning or adopting a more robust model architecture could help optimize classification accuracy and reduce misclassification rates.